

# ICAPS Summer School 2020: Probabilistic Planning (MDPs)

**Scott Sanner**



UNIVERSITY OF  
**TORONTO**



# Model-based Probabilistic Planning

Mausam

Computer Science and Engineering

Indian Institute of Technology (IIT) Delhi

# Planning under Uncertainty


- Definition:

Computing sequences of actions to obtain occasional rewards in a known, stochastic environment

# Applications

# Elevator Control

- **Concurrent Actions**
  - Elevator: up/down/stay
  - 6 elevators:  $3^6$  actions
- **Dynamics:**
  - Random arrivals (e.g., Poisson)
- **Objective:**
  - Minimize total wait
  - (Requires being proactive about future arrivals)
- **Constraints:**
  - People might get annoyed if elevator reverses direction



# Two-player Games

- **Othello / Reversi**

- Solved by Logistello!
- Monte Carlo RL (self-play)  
+ Logistic regression + **Search**



- **Backgammon**

- Solved by TD-Gammon!
- Temporal Difference (self-play)  
+ Artificial Neural Net + **Search**



- **Go**

- Learning + **Search**
- AlphaGo (MCTS + deep learning)  
recently the world champion



# Multi-player Games: Poker

- **Multiagent (adversarial)**
  - Opponent may abruptly change strategy
  - Might prefer best outcome for *any* opponent strategy (e.g, a Nash equilibrium)
- **Multiple rounds (sequential)**
- **Partially observable!**
  - Earlier actions may reveal information
  - Or they may not (bluff)





# DARPA Grand Challenge

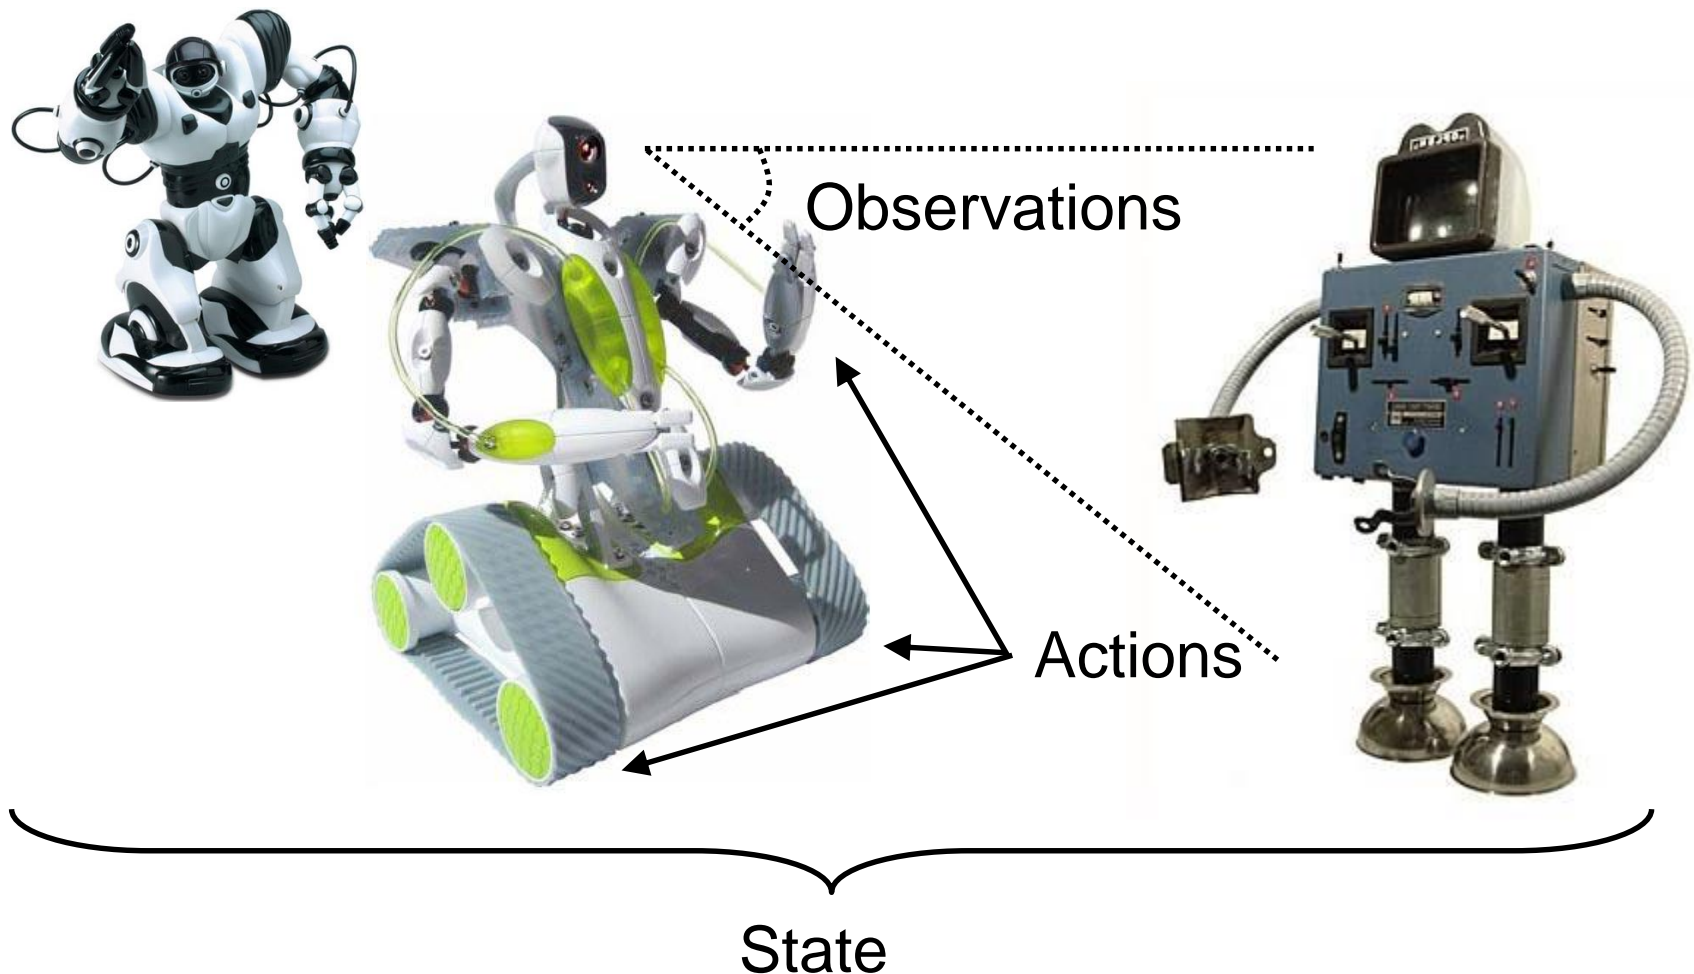
- **Autonomous mobile robotics**
  - Extremely complex task, requires expertise in vision, sensors, real-time operating systems
- **Partially observable**
  - e.g., only get noisy sensor readings
- **Model unknown**
  - e.g., steering response in different terrain





**How to model  
these problems?**

# Observations, States, & Actions



# Observations

- **Observation set  $O$** 
  - Perceptions, e.g.,
    - Distance from car to edge of road
    - My opponent's bet in Poker

# States

- **State set  $S$** 
  - At any point in time, system is in some state
    - Actual distance to edge of road
    - My opponent's hand of cards in Poker

# Agent Actions

- **Action set  $A$** 
  - Actions could be *concurrent*
  - If  $k$  actions,  $A = A_1 \times \dots \times A_k$ 
    - Schedule all deliveries to be made at 10am

# Agent Actions

- **Action set A**

- All actions need not be under agent control

- Other agents, e.g.,

- Alternating turns: Poker, Othello

- Concurrent turns: Highway Driving, Soccer

- *Exogenous events* due to *Nature*, e.g.,

- Random arrival of person waiting for elevator

- Random failure of equipment

- If uncontrolled, model as random variables

# Observation Function

- How to relate states and observations?
  - ***Not observable:***
    - $O = \emptyset$
    - e.g., heaven vs. hell
      - » only get feedback once you meet St. Pete
  - ***Fully observable:***
    - $S \leftrightarrow O$  ... the case we focus on!
    - e.g., many board games,
      - » Othello, Backgammon, Go
  - ***Partially observable:***
    - all remaining cases
    - e.g., driving a car, Poker, the real world!



# Recap

- So far
  - Actions
  - States
  - Observations
- How to map between
  - Previous states, actions, and future states?
  - States and observations?
  - States, actions and rewards?
  - Sequences of rewards and optimization criteria?

# Transition Function

- How do actions take us between states?
  - $T(s,a,s')$  encodes  $P(s'|s,a)$
  - Some properties
    - *Stationary*: T does not change over time
    - *Markovian*: Only depends on previous state / action
    - If T not Markovian or stationary
      - can sometimes achieve by augmenting state description
        - » e.g., elevator traffic differs throughout day...
        - encode time in state to make T Markovian!


# Goals and Rewards

- Goal-oriented rewards
  - Assign any reward value s.t.  $R(\text{success}) > R(\text{fail})$
  - Can have negative costs  $C(a)$  for action  $a$
- What if multiple (or no) goals?
  - How to specify preferences?
  - $R(s,a)$  assigns utilities to each state  $s$  and action  $a$ 
    - Then *maximize expected reward (utility)*



But, how to trade off  
rewards over time?

# Optimization: Best Action when $s=1$ ?



- Must define objective criterion to optimize!
  - How to trade off immediate vs. future reward?
  - E.g., use discount factor  $\gamma$  (try  $\gamma=.9$  vs.  $\gamma=.1$ )

# Trading Off Sequential Rewards

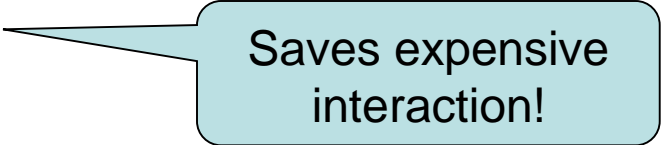
- Sequential-decision making objective
  - Horizon
    - *Finite*: Only care about h-steps into future
    - *Infinite*: Literally; will act same today as tomorrow
  - How to trade off reward over time?
    - *Expected average cumulative return*
    - *Expected discounted cumulative return*
      - Use discount factor  $\gamma$
      - Reward t time steps in future discounted by  $\gamma^t$

# Recap

- Model so far
  - Actions  $A$
  - States  $S$
  - Observation  $O$
  - Transition function  $T$ :  $P(s'|s,a)$
  - Observation function  $Z$ :  $P(o'|s,a)$  – *POMDPs only*
  - Reward function:  $R(s,a)$
  - Optimization criteria
- But are the above
  - Known or unknown?

# Knowledge of Environment

- **Model-known:**
  - Know observation, transition, & reward functions
  - Called: *Planning (under uncertainty)*
    - Planning generally assumed to be goal-oriented
    - *Decision-theoretic* if maximizing expected utility
- **Model-free:**
  - $\geq 1$  unknown: observation, transition, & reward functions
  - Called: *Reinforcement learning*
    - Have to interact with environment to obtain samples
- **Model-based: approximate model in model-free case**
  - Permits hybrid planning and learning



Saves expensive interaction!




# Finally a Formal Model

- **Define the previous model**
  - MDP:  $\langle \mathbf{S}, \mathbf{A}, \mathbf{T}, \mathbf{R} \rangle$
  - POMDP:  $\langle \mathbf{S}, \mathbf{A}, \mathbf{O}, \mathbf{Z}, \mathbf{T}, \mathbf{R} \rangle$
  - Whether known / unknown
- **Characterize the solutions**
  - And efficiently find them!

# **Model-based Solutions to MDPs**

# MDPs $\langle S, A, T, R \rangle$


Note: fully observable



- $\mathbf{S} = \{1,2\}$ ;  $\mathbf{A} = \{\text{stay, change}\}$
- **Reward**
  - $R(s=1, a=\text{stay}) = 2$
  - ...
- **Transitions**
  - $T(s=1, a=\text{stay}, s'=1) = P(s'=1 \mid s=1, a=\text{stay}) = .9$
  - ...

How to act  
in an MDP?  
  
Define policy  
 $\pi: \mathbf{S} \rightarrow \mathbf{A}$

# What's the best Policy?





- Must define reward criterion to optimize!
  - Discount factor  $\gamma$  important ( $\gamma=1.0$  vs.  $\gamma=0.1$ )

# Between Value and Policy Iteration

- *Value iteration*
  - Each iteration seen as doing 1-step of policy evaluation for current greedy policy
  - Bootstrap with value estimate of previous policy
- *Policy iteration*
  - Each iteration is full evaluation of  $V_{\pi}$  for current policy  $\pi$
  - Then do greedy policy update
- *Modified policy iteration*
  - Like policy iteration, but  $V_{\pi_i}$  need only be closer to  $V^*$  than  $V_{\pi_{i-1}}$ 
    - Fixed number of steps of successive approximation for  $V_{\pi_i}$  suffices when bootstrapped with  $V_{\pi_{i-1}}$
  - Typically faster than VI & PI in practice


# Canonical Example: Grid World

- The agent lives in a grid
- Walls block the agent's path
- The agent's actions do not always go as planned:
  - 80% of the time, the action North takes the agent North (if there is no wall there)
  - 10% of the time, North takes the agent West; 10% East
  - If there is a wall in the direction the agent would have been taken, the agent stays put
- Big rewards come at the end



# Solving MDPs

- In an MDP, we want an optimal **policy**  $\pi^*: S \times 0:H \rightarrow A$ 
  - A policy  $\pi$  gives an action for each state for each time



- An optimal policy maximizes expected sum of rewards
- Contrast: In deterministic, want an optimal **plan**, or sequence of actions, from start to a goal



# Outline

- Optimal Control

=

given an MDP ( $S, A, T, R, \gamma, H$ )

find the optimal policy  $\pi^*$

- Exact Methods:

- ***Value Iteration***

- Policy Iteration

# Value Iteration

- Algorithm:

- Start with  $V_0^*(s) = 0$  for all  $s$ .

- For  $i=1, \dots, H$

Given  $V_i^*$ , calculate for all states  $s \in S$ :

$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + V_i^*(s')]$$

- This is called a **value update** or **Bellman update/back-up**

- $V_i^*(s)$  = the expected sum of rewards accumulated when starting from state  $s$  and acting optimally for a horizon of  $i$  steps

# Value Iteration

```
values = {each state : 0}
loop ITERATIONS times:
  previous = copy of values
  for all states:
    EVs = {each legal action : 0}
    for all legal actions:
      for each possible next_state:
        EVs[action] += prob * previous[next_state]
  values[state] = reward(state) + discount * max(EVs)
```

# Value Iteration in Gridworld

noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$

0.00 ▶	0.00 ▶	0.00 ▶	1.00
0.00 ▶		◀ 0.00	-1.00
0.00 ▶	0.00 ▶	0.00 ▶	0.00 ▼

VALUES AFTER 1 ITERATIONS

# Value Iteration in Gridworld

noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$

0.00 ▶	0.00 ▶	0.72 ▶	1.00
0.00 ▶		▲ 0.00	-1.00
0.00 ▶	0.00 ▶	0.00 ▶	0.00 ▼

$$0.72 = 1 * .8 * .9$$

VALUES AFTER 2 ITERATIONS

# Value Iteration in Gridworld

noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$


0.00 ▶	0.52 ▶	0.78 ▶	1.00
0.00 ▶		▲ 0.43	◻ -1.00
0.00 ▶	0.00 ▶	▲ 0.00	0.00 ▼

VALUES AFTER 3 ITERATIONS

$$0.78 = 0.72 + 1 * .9 * .8 * .9$$

# Value Iteration in Gridworld

noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$




**VALUES AFTER 4 ITERATIONS**



# Value Iteration in Gridworld


noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$



VALUES AFTER 5 ITERATIONS

# Value Iteration in Gridworld


noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$



VALUES AFTER 100 ITERATIONS


# Value Iteration in Gridworld

noise = 0.2,  $\gamma = 0.9$ , two terminal states with  $R = +1$  and  $-1$



**VALUES AFTER 1000 ITERATIONS**

# Exercise 1: Effect of discount, noise



- (a) Prefer the close exit (+1), risking the cliff (-10)
- (b) Prefer the close exit (+1), but avoiding the cliff (-10)
- (c) Prefer the distant exit (+10), risking the cliff (-10)
- (d) Prefer the distant exit (+10), avoiding the cliff (-10)

- (1)  $\gamma = 0.1$ , noise = 0.5
- (2)  $\gamma = 0.99$ , noise = 0
- (3)  $\gamma = 0.99$ , noise = 0.5
- (4)  $\gamma = 0.1$ , noise = 0

# Exercise 1 Solution

0.00 ▸	0.00 ▸	0.01 ▼	0.01 ▸	0.10 ▼
0.00 ▼		0.10 ▼	0.10 ▸	1.00 ▼
0.00 ▼		1.00		10.00
0.00 ▸	0.01 ▸	0.10 ▲	0.10 ▸	1.00 ▲
-10.00	-10.00	-10.00	-10.00	-10.00

(a) Prefer close exit (+1), risking the cliff (-10) ---  $\gamma = 0.1$ , noise = 0

# Exercise 1 Solution

0.00 ▶	0.00 ▶	0.00 ▼	0.00 ▼	0.03 ▼
▲ 0.00		0.05 ▼	0.03 ▶	0.51 ▼
0.00 ▼		1.00		10.00
▲ 0.00	▲ 0.00	▲ 0.05	▲ 0.01	▲ 0.51
-10.00	-10.00	-10.00	-10.00	-10.00

(b) Prefer close exit (+1), avoiding the cliff (-10) --  $\gamma = 0.1$ , noise = 0.5

# Exercise 1 Solution

9.41 ▶	9.51 ▶	9.61 ▶	9.70 ▶	9.80 ▼
9.32 ▼		9.70 ▶	9.80 ▶	9.90 ▼
9.41 ▼		1.00		10.00
9.51 ▶	9.61 ▶	9.70 ▶	9.80 ▶	9.90 ▲
-10.00	-10.00	-10.00	-10.00	-10.00

(c) Prefer distant exit (+1), risking the cliff (-10) --  $\gamma = 0.99$ , noise = 0

# Exercise 1 Solution



(d) Prefer distant exit (+1), avoid the cliff (-10) --  $\gamma = 0.99$ , noise = 0.5



# Value Iteration Convergence

**Theorem.** Value iteration converges. At convergence, we have found the optimal value function  $V^*$  for the discounted infinite horizon problem, which satisfies the Bellman equations

$$\forall S \in \mathcal{S} : \quad V^*(s) = \max_A \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Now we know how to act for infinite horizon with discounted rewards!

- Run value iteration till convergence.
- This produces  $V^*$ , which in turn tells us how to act, namely following:

$$\pi^*(s) = \arg \max_{a \in A} \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$$

- Note: the infinite horizon optimal policy is stationary, i.e., the optimal action at a state  $s$  is the same action at all times. (Efficient to store!)

# Convergence and Contractions

- Define the max-norm:  $\|U\| = \max_s |U(s)|$

- Theorem: For any two approximations U and V

$$\|U_{i+1} - V_{i+1}\| \leq \gamma \|U_i - V_i\|$$

- I.e. any distinct approximations must get closer to each other, so, in particular, any approximation must get closer to the true U and value iteration converges to a unique, stable, optimal solution

- Theorem:

$$\|V_{i+1} - V_i\| < \epsilon, \Rightarrow \|V_{i+1} - V^*\| < 2\epsilon\gamma/(1 - \gamma)$$

- I.e. once the change in our approximation is small, it must also be close to correct

# Policy Evaluation

- Recall value iteration iterates:

$$V_{i+1}^*(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_i^*(s')]$$

- Policy evaluation:

$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$

- At convergence:

$$\forall s \quad V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$

# Policy Iteration

- Alternative approach:
  - **Step 1: Policy evaluation:** calculate utilities for some fixed policy (not optimal utilities!) until convergence
  - **Step 2: Policy improvement:** update policy using one-step look-ahead with resulting converged (but not optimal!) utilities as future values
  - Repeat steps until policy converges
- This is **policy iteration**
  - It's still optimal!
  - Can converge faster under some conditions

# Policy Evaluation Revisited

- Idea 1: modify Bellman updates

$$V_0^\pi(s) = 0$$

$$V_{i+1}^\pi(s) \leftarrow \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_i^\pi(s')]$$


- Idea 2: it's just a linear system, solve with Matlab (or whatever),  
variables:  $V^\pi(s)$ ,  
constants:  $T, R$

$$\forall s \quad V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$


# Advanced (PO)MDP Modeling with RDDDL

# What is RDDDL?

- Relational Dynamic Influence Diagram Language
  - Relational [DBN + Influence Diagram]
  - Everything is a fluent!
    - states
    - observations
    - actions
  - Conditional distributions are probabilistic programs



# Wildfire Domain



- Contributed by Zhenyu Yu (School of Economics and Management, Tongji University)
  - Karafyllidis, I., & Thanailakis, A. (1997). *A model for predicting forest fire spreading using gridular automata*. Ecological Modelling, 99(1), 87-97.



# Wildfire in RDDDL

```
cpfs {

    burning'(?x, ?y) =
        if ( put-out(?x, ?y) )
            then false
        else if (~out-of-fuel(?x, ?y) ^ ~burning(?x, ?y))
            then Bernoulli( 1.0 / (1.0 + exp[4.5 - (sum_{?x2: x_pos, ?y2: y_pos}
                (NEIGHBOR(?x, ?y, ?x2, ?y2) ^ burning(?x2, ?y2)))])) )
        else
            burning(?x, ?y); // State persists

    out-of-fuel'(?x, ?y) = out-of-fuel(?x, ?y) | burning(?x,?y);

};

reward =
    [sum_{?x: x_pos, ?y: y_pos} [ COST_CUTOUT*cut-out(?x, ?y) ]]
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_PUTOUT*put-out(?x, ?y) ]]
+ [sum_{?x: x_pos, ?y: y_pos} [ COST_NONTARGET_BURN*[ burning(?x, ?y) ^ ~TARGET(?x, ?y) ]]]
+ [sum_{?x: x_pos, ?y: y_pos}
    [ COST_TARGET_BURN*[ (burning(?x, ?y) | out-of-fuel(?x, ?y)) ^ TARGET(?x, ?y) ]]];
```

# Facilitating Model Development by Writing Simulators: Relational Dynamic Influence Diagram Language (RDDL)

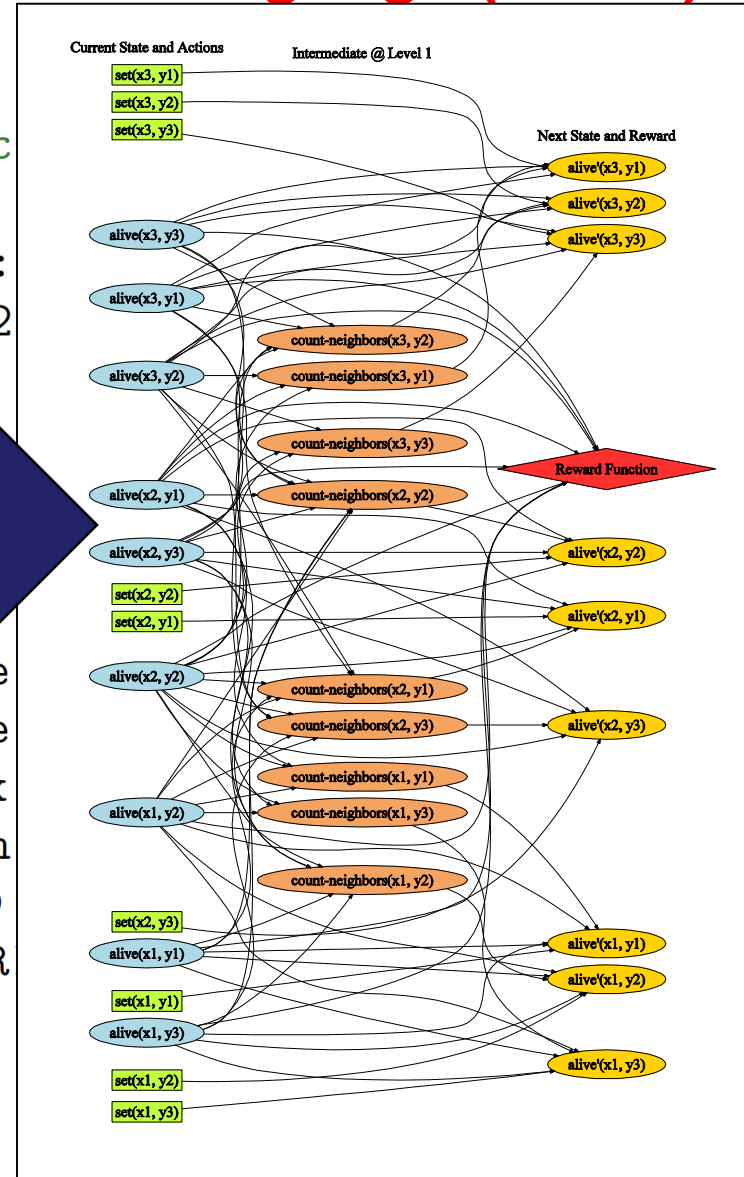
```
// Store alive-neighbor count for each
count-neighbors(?x,?y) =
  KronDelta(sum_{?x2 : x_pos, ?y2 :
    [NEIGHBOR(?x,?y,?x2,?y2
```

```
// Determine whether
alive'(?x,?y) = if (
else
```

**Automatic  
Translation**

Write  
probabilistic  
programs for  
transitions

```
^ (count-ne
^ (count-ne
| [~alive(?x
^ (count-n
| set(?x,?y)
then Bernoulli(PROB_R
else Bernoulli(1.0 -
```



# RDDLSim Software

Open source & online at

<http://code.google.com/p/rddlsim/>

# RDDL Software Overview

- BNF grammar and parser
- Simulator
- Automatic compilation / translations
  - LISP-like format (easier to parse)
  - SPUDD & Symbolic Perseus (boolean subset)
  - Ground PPDDL (boolean subset)
- Client / Server
  - Java and C/C++ sample clients
  - Evaluation scripts for log files
- Visualization
  - DBN Visualization
  - Domain Visualization – see how your planner is doing

# Initial Use of RDDDL

- Have run two major competitions at ICAPS
- Translations to draw in different communities
  - UAI Factored MDP / POMDP community
  - ICAPS PPDDL community
  - 11 competitors in 2011, 6 competitors in 2014
- Competitions drive research progress!
  - Historically, ICAPS focused on deterministic replanning
  - With RDDDL + new domains, **MCTS dominates**  
(namely PROST system by Thomas Keller *et al*)